

First Person Shooter level generation using Generative Adversarial Networks

Edoardo Giacomello

15 febbraio 2018

Abstract

Estratto in lingua Italiana

To someone...

Acknowledgments

Contents

Abstract	3
Estratto in lingua Italiana	5
Acknowledgments	9
1 Introduction	17
1.1 Background: Level Design	17
1.2 Related Work	17
1.2.1 Procedurally Generated Content	17
1.2.2 Procedural Content Generation via Machine Learning (PCGML)	17
1.2.3 Video Game Level Corpus	17
1.2.4 Generative Adversarial Networks	17
1.3 Scope	17
1.4 Thesis Structure	17
1.5 Summary	17
2 Towards learn-based level generation	19
2.0.1 Overview	19
2.1 Generative Adversarial Networks	19
2.1.1 Overview	19
2.1.2 Deep Convolutional GAN	20
2.1.3 Wesserstein GAN	20
2.1.4 Wesserstein GAN with Gradient Penalty	20
2.2 Game of choice: DOOM	20
2.2.1 Description	20
2.2.2 Motivation	20
2.3 Summary	20
3 Dataset and Data Representation	21
3.1 Data Sources	21
3.2 Source Data Format: WAD Files	22
3.2.1 Overview	22
3.2.2 Doom Level Format	23
3.2.3 Conversion issues	24
3.3 Target Data Format: Feature Maps and Vectors	26
3.3.1 Level Description and Motivation	26
3.3.2 Feature Maps	26
3.3.3 Graph Representation	27
3.3.4 Text Representation	28
3.3.5 Scalar Features	28
3.4 Dataset Organization	28
3.4.1 Overview	28

3.4.2	Full Dataset and Filtered Dataset	28
3.4.3	Level Size Statistics	30
3.5	Summary	31
4	System Design and Overview	43
4.1	Generative Model Structure	43
4.2	Use Cases	44
4.2.1	Use Case: Model Optimization (Training)	45
4.2.2	Use Case: Model Validation and Sample Evaluation	45
4.2.3	Use Case: Sampling or Generation	46
4.3	Data Flow	46
4.4	WAD Editor and Feature Extractor	48
4.5	Summary	48
5	System Architecture	49
5.1	Component View	49
5.2	Neural Network Architecture	49
6	Experiment Design and Results	51
6.1	Neural Network Architecture and Training Algorithm	51
6.1.1	Network Architecture and Hyper-parameters	51
6.1.2	Techniques and "GAN Tricks" used	51
6.1.3	Training Algorithm	51
6.1.4	Input data	51
6.1.5	Resulting Model	51
6.1.6	Evaluation metrics	51
6.2	Sampling the network	51
6.3	Generated Samples	51
6.4	In-Game Demonstration	51
6.5	Summary	51
7	Results Evaluation and Conclusions	53
7.1	Results Evaluation	53
7.1.1	Samples Evaluation	53
7.1.2	Loss of accuracy	53
7.2	Summary	53
8	Future Work	55
8.1	Open Problems	55
8.2	Possible Applications and future develops	55
9	Appendix	57
	Glossary	61

List of Figures

3.1	A simple level showing sectors and linedefs	25
3.2	DoomDataset Size distribution	30
4.1	System Overview: Generative Model Structure	44
4.2	Use Case: Model Optimization	45
4.3	Use Case: Validation and Sample Evaluation	46
4.4	Use Case: Sampling or Generation	47
4.5	System Overview: Data Flow Diagram	47

List of Tables

3.1	WAD File structure	32
3.2	ThingsMap Encoding (1 of 3)	33
3.3	ThingsMap Encoding (2 of 3)	34
3.4	ThingsMap Encoding (3 of 3)	35
3.5	TriggerMap Encoding	36
3.6	Textual Representation Encoding	36
3.7	Features: IDArchive Metadata	37
3.8	Features: WAD-extracted	38
3.9	Features: PNG-extracted (1 of 2)	39
3.10	Features: PNG-extracted (2 of 2)	40
3.11	Features: Graph	41
3.12	Percentiles of level width and height distributions	41

Chapter 1

Introduction

1.1 Background: Level Design

1.2 Related Work

1.2.1 Procedurally Generated Content

1.2.2 Procedural Content Generation via Machine Learning (PCGML)

1.2.3 Video Game Level Corpus

1.2.4 Generative Adversarial Networks

1.3 Scope

1.4 Thesis Structure

1.5 Summary

Chapter 2

Towards learn-based level generation

2.0.1 Overview

2.1 Generative Adversarial Networks

2.1.1 Overview

GAN A model that in the latest years gained increasingly more interest in the field of Machine Learning research is the one of the Generative Adversarial Network, proposed by Goodfellow et al. in “Generative Adversarial Networks” (2014) [12]. The main idea of this kind of generative model is to use two neural networks which are posed in an adversarial setting that models a two-player Minimax Game [8, p. 276]. In particular, a generative network called G is trained to capture the data distribution while a discriminator network D estimates the probability that each sample comes from the real data distribution rather than the one generated by G.

Conditional GAN In our work we make use of a modification of the GAN model introduced by Mirza and Osindero [17], which allows to condition the data generation to class labels. This structure is summarized in section 4.1 where the logical structure of the generative model we adopt is presented.

Additional Architecture Due to the amount of interest that GANs have raised, many researches have been conducted to improve the quality performances of GANs, leading to a variety of proposed new architectures. They can be classified as modifications to the underlying neural network architectures or even to the conceptual setting of the model itself, by means of changes to the loss function or the training algorithm. Other, less formal but still effective adjustment to the proposed models are the so called ”tricks” in machine learning discussion communities, which adoption is often suggested in order to improve the difficult training of the network. The final architecture adopted is described in section 6.1.1.

2.1.2 Deep Convolutional GAN

2.1.3 Wesserstein GAN

2.1.4 Wesserstein GAN with Gradient Penalty

2.2 Game of choice: DOOM

2.2.1 Description

2.2.2 Motivation

2.3 Summary

Chapter 3

Dataset and Data Representation

Overview This chapter aims to be a description of the dataset structure the model is trained and evaluated with, as well as an overview of the processes that led to the creation of the dataset themselves. In section 3.1 a reference to the data sources is given, then the focus of section 3.2 is on how data is natively encoded for the game engine in order to give some hints on what are the difficulties to face in converting to and from that format in an automatic way. Section 3.3 describes in detail what data is provided with the dataset, how levels are converted from the native format and what features are extracted in order to provide an input for the neural network. Lastly, section 3.4 gives an overview of how the data is provided and organized in the resulting dataset.

3.1 Data Sources

Archive All data used to train and validate the model comes from the *Idgames Archive* founded in 1994 by Barry Bloom [2] and mirrored on various FTP sources. The mirror we used for collecting levels is Doomworld.com [6], which is one of the oldest and currently most active community about the DOOM video-game series [5].

Level Selection Idgames archive includes levels for multiple games such as DOOM, DOOM 2 or their various modifications. They are divided in hierarchical categories, which classify levels by game, game mode (multi-player "deathmatch" or single-player), and alphabetically. Amongst all the categories we selected only those levels that belong to "DOOM" and "DOOM 2", excluding sub-categories named "deathmatch" and "Ports". This choice has been made in order to avoid mixing possibly different kinds of levels, since a level designed for a Single Player Mode could be structurally different from a level which is designed for multi-player games. Moreover the "Ports" category has been excluded because levels contained in it are intended to work with modifications of the game engine code, and it would have led to problems in managing every particular exceptional behaviour.

Source Data Organization Levels in Idgames archive are stored in zip archives, including a "READ ME" text file containing author notes and the WAD file that contains up to 32 levels. Each zip archive can be downloaded from the respective download page, which presents a variable quantity of information such as: author, a short description, screen shots, user reviews, number of views and downloads, etc. The dataset we present in this work always keeps track of these information about each level for correct attribution. It also offers a "snapshot" of average user review score and the number of views and downloads, when available. It is worth noting that since Doomworld website recently switched to a different download system [5], data may not always be accurate, especially that concerning download and view counts; they are still proposed as a starting point for further research.

3.2 Source Data Format: WAD Files

Introduction The Doom Game Engine [14] makes use of package files called "WADs" to store every game resource such as Levels, Textures, Sounds, etc. WAD files have been designed in order to make the game more extendible and customizable, and opened the way for a considerably large amount of user-generated content. This section is not meant to be a complete description of how WADs files are structured, but only an overview of which aspects we considered for writing the software that generates the dataset from the WAD files. Every information about WADs files has been taken from the *The Unofficial Doom Specs* [9] and we refer to that document a deeper explanation on every aspect of the file format.

3.2.1 Overview

Type of WADs There are two types of WADs: the "IWAD", or "Internal WAD", and the "PWAD" or "Patch WAD". The original game files, called "DOOM.WAD" and "DOOM2.WAD", are of the "IWAD" type, as they contain every asset that is needed for the game to run. The WADs containing custom content or modifications to existing content are of the "PWAD" type. The content which is defined in a PWAD is added or replaced to the original IWAD when the WAD is loaded. For example, if a PWAD defines the level "MAP01", which is already defined in "DOOM2.WAD", the PWAD level is loaded instead of the original one, while maintaining all the other content unaltered. Since in our work we deal only with PWADs, we will generally refer to them simply as WADs.

Lumps Every data inside a WAD is stored as a record called Lump, which has a name up to 8 characters and a structure and size that is different depending on the lump type. Generally there are no restrictions on lump order with the exception of some of them, including those needed for defining a Level.

WAD Structure Every WAD file is divided in three sections: A header, a set of Lumps and a trailing Directory. The header holds the information about the WAD type, the number of Lumps and the location of the Directory, which is positioned after the last Lump. The directory contains one 12-bytes entry for each Lumps that specifies the Lumps location, size and name. Table 3.1 reports a simplified description of a WAD file.

Coordinate Units The Doom game engine describes coordinates using integer values between -32768 and +32767, and it is proportional to one pixel of a texture. This unit is called "Map Unit" or "Doom Unit" in this work. Although there's not an unique real world interpretation of one Map Unit, we used an approximation for which each relevant map tile or image pixel is 32x32 MU large; this choice is motivated by the fact that the smallest radius of a functional object in DOOM is 16 MU.

3.2.2 Doom Level Format

Overview Level data in a WAD file follows a precise structure. In particular each level is composed of an ordered sequence of lumps that describes its structure:

(NAME) : Name of the level slot in DOOM or DOOM2 format.

THINGS List of every game object ("Thing") that is placeable inside the level.

LINEDEFS List of every line that connects two vertices.

SIDEDEFS A list of structures describing the sides of every Linedef.

VERTEXES Unordered list of vertices.

SEGS A list of linedef segments that forms sub-sectors.

SSECTORS A list of sub-sectors, which are convex shapes forming sectors.

NODES A binary tree sorting sub-sectors for speeding up the rendering process.

SECTORS A list of Sectors. A Sector is a closed area that has the same floor and ceiling height and textures.

REJECT Optional lump that specifies which sectors are visible from the other. Used to optimize the AI routines.

BLOCKMAP Pre-computed collision detection map.

Essential Subset and Optimizations It is important to notice that although all the lumps above (with the exception of REJECT) are mandatory to build a playable level, some of them can be automatically generated from the remaining ones using external tools. In particular, an editor software or designer has to provide at least the lumps (name), THINGS, LINEDEFS, SIDEDEFS, VERTEXES and SECTORS. The lumps SEGS, SSECTORS, NODES, REJECT and BLOCKMAP serve the purpose of speeding-up the rendering process by avoiding runtime computation. In particular the Doom Engine uses a Binary Space Partitioning Algorithm [10] for pre-computing the Hidden Surface Determination (or occlusion culling), and it is usually done by an external tool. In this work we used the tool "*BSP v5.2*" [4] in the last stage of the pipeline, in order to produce playable DOOM levels from the network output. In the following paragraphs only the lump types that are not generated by the external tool are described.

(Name) The first lump of a level is its slot name. We indicate this lump between parenthesis because, differently from the other lumps, this one has no data associated. The Name field in the Directory is the slot name itself and the size is therefore zero. The level name descriptor has to match either *ExMy* ("Episode x, Map y") or *MAPzz* for DOOM or DOOM2 respectively, where x ranges from 1 to 4, y from 1 to 9 and zz from 1 to 32.

Things A doom "Thing" is every object included in a level that is not a wall, pavement, or a door. A "Things" lump is a list of entries each one containing five integers specifying the *position* (x,y) in MU, the *angle* the thing is facing, the *Thing type* index, and a set of flags indicating in which difficulty level the thing is present and whether the thing is deaf, if it is an enemy.

Linedefs A Linedef is any line that connects two vertices. Since DOOM maps are actually bi-dimensional, a single line is needed to define each wall or step. Linedefs do not necessarily have to match with visible entities, as each Linedef can also represent invisible boundaries or triggers, which can be thought as tripwires or switches that make something happen to a sector. If the Linedef acts as a trigger than it references another sector and has a type which defines what would happen to it when the Linedef is activated. For example, a door is implemented as a sector that raises up to the ceiling when a certain Linedef is triggered, but this behaviour is specified in the switch and not in the door as one might think. All the options are defined by a set of flags that specifies what kind of objects the Linedef blocks, if it's two sided or not, the trigger activation condition and so on. Finally, each Linedef has to specify what "vertical plane" (or Sidedef) lies on its right and left side. Only the "left sidedef" field can be left empty, due to the way the sectors are represented.

Sidedefs A Sidedef is a structure that contains the texture data for each linedef. It corresponds to the side of each wall and it's referenced by the linedefs. Other than the options for texture visualization, it also specifies the sector number that this plane is facing, implicitly defining the sector boundaries.

Vertexes ¹ This is the simplest type of Lump, consisting in an unordered list of map coordinates in MU. Each entry has an x and y position expressed in a 16-bit signed integer. Linedefs references the starting and ending vertex by reporting their index in this list.

Sectors A sector is defined as any area that have constant floor and ceiling heights and textures. This definition highlights the fact that the doom engine is in fact a 2d engine, since it's not possible to define a sector above or below another one. A sector does not necessarily have to be closed nor a single connected polygon, but non-closed sectors can cause some issues during gameplay. The only constraint to define sectors comes from the fact that linedefs must have a right sidedef, but the left one is optional; the sectors are normally described as (a set of) positively oriented curves [7], with the exception of linedefs that are shared with another sector that may be reversed. The fields of a sector lump includes the floor and ceiling height, the name of the floor and ceiling texture, the light level, the "type" to controls some lighting effects and damaging floor, and the tag number that is referenced by the linedef triggers. Figure 3.1 shows an example of a level with 3 sectors.

3.2.3 Conversion issues

Since 1994 many DOOM players started producing a large amount of levels and increasingly sophisticated editors came to light. This led to a notable variety in conventions, optimizations and use of bad practices that we tried to deal with in developing the Python module for reading and writing wad files. Some of these practices includes unnecessary lumps, random-data instead of null padding for names, and other inconsistencies or arbitrary conventions. However, we paid particular attention during the writing phase in order to precisely follow the specifications and avoid generating low quality WAD files as much as we could. Some other difficulties came with the necessity to render wad files as images and vice-versa. The first one is that neither Linedefs nor vertices came in an any ordered format, along with the fact that sector vertices are not explicitly defined, but only referenced through the linedef/sidedef chain. This means that for finding a sector shape one has to find all sidedefs with the desired sector number, then find all the linedefs referencing those sidedefs and lastly retrieving the vertices, leading to an increase of complexity. Another problem was the fact that although sectors must be positively oriented curves, it is not mandatory to use the left sidedefs for adjacent sectors, leading to duplicated linedefs in the opposite direction. Even some optimizations such as sidedef compression may be possible, that is referencing the same sidedef wherever the same wall texture is used, adding

¹Although the correct spelling should be "vertices", we keep the original version of the field name.




Figure 3.1: A simple level showing three sectors A, B and C and the linedefs defining them, following the positive oriented curve constraint. Sector C can be viewed as a small platform inside the sector B. Solid arrows represent walls, while dashed lines represent invisible linedefs or changes in height between two sectors (steps). In this level, every solid arrow is a linedef that specifies only a right sidedef, with the exception of the one separating sectors A and B that has both a right and a left one.

complexity to the conversion script. Moreover, the concept of sector and the concept of room are not equivalent: even though a sector is defined as an area of constant height, this does not enforce to define a sector only where height changes, so the semantic of a sector actually depends on the designer and the editor used.

3.3 Target Data Format: Feature Maps and Vectors

Introduction This section provides a description of the data format as it is stored in the level dataset.

3.3.1 Level Description and Motivation

Levels are read as a structured object by a module we wrote for the purpose of providing developers and designers a programmatic way to access, analyse and edit DOOM WAD files, instead of using visual authoring software. This allows to automatise some tasks that would be very long to accomplish with standard editors or tools and also offers developers the possibility to write custom editor and scripts.

In order to provide the most complete set of information possible every level is converted to a set of images, called Feature Maps in this work, a tiled representation in textual format that is an extension of the one taken from [20], a graph representation and a set of textual and scalar features that contains both the WAD meta-data and the level features and metrics calculated either on the WAD representation (sectors, subsectors, linedefs, etc), the Feature Map representation or the graph representation of the level. A detailed list and explanation of each feature is provided in section 3.3.5.

The choice of these representations have been made primarily for the need of having a data format that the generator model can work on. In particular, convolutional neural networks are naturally designed to work well with bi-dimensional or tri-dimensional data such as (multi-channel) images, for this reason the image representation arose naturally. The text representation is provided mainly for consistency with the data format given in [20], even if our representation uses two characters per tile instead of one. Finally, scalar and graph representation have been collected for the need of quantifying and summarising some properties and having a more abstract representation of a level, which can be very helpful in the case this data had to be used in other works.

3.3.2 Feature Maps

Feature Maps are a set of images each of them describing a different aspect of the level. In particular we used for each Feature Map a grayscale 8-bit image in which each pixel can assume values between 0 and 255. This allowed to obtain a good degree of precision while still maintaining a reasonable dataset size.

Because of the motivations explained in section 3.2.1, each pixel in a Feature Map corresponds to a square of 32x32 MU. In the following paragraph we will describe in detail each of them along with the data encoding.

FloorMap The FloorMap is the most basic form of level representation, since it only represents which part of the space are occupied by the level and which are empty. This kind of map is often used in robotic mapping.

This map also describes approximatively the level area that is possible to traverse, since in DOOM walls have no thickness.

"Floor" pixels have value 255 (white) and "Empty" pixels have value 0 (black).

WallMap The WallMap represent the impassable walls of the level. They are represented as a one-pixel-wide line and are obtained by directly drawing each Linedef with the impassable flag

set on a black image.

Pixel values are 0 for Empty area or floors and 255 for the walls.

HeightMap The HeightMap is another common map used for visualizing the height of a certain surface. Since height level in DOOM levels are completely arbitrary and virtually unbounded, we normalize each level between its lowest and highest height, assigning the pixel value 0 for empty parts of the level and the remaining are calculated from the formula $c_h = \lfloor h * \frac{255}{|H|} \rfloor$,

where H is the ordered set of possible height values for the level and $h \in \{1, 2, \dots, |H|\}$ is the index of the height value in H for which we want to calculate the encoded colour.

For example if a level takes the height values $H = \{0, 10, 15, 20\}$ they will be encoded respectively as $c_h = \{63, 127, 191, 255\}$ and 0 for the empty areas. Although this map loses the information about the differences between a height level and another, it has the advantage to represent "higher" and "lower" parts of the level without polarizing the entire map due to extreme levels: while the majority of the levels has a few changes that can be approximated as uniform (such as levels with stairs connecting a few rooms), other had some extreme changes in height but only for a small portion of the map (like a very high elevator leading to a small secret room) that led to scaling problems.

ThingsMap The ThingsMap represent only data that is contained in the THINGS lump. It features a series of pixels placed at the thing coordinate, with a value that corresponds to a particular "thing". Pixel colours have been grouped by functional purposes so for example weapons occupies values that are close each other. This is for tolerating some output noise during generation without completely changing the functional aspect of an object as would have happened if we kept the original things encoding. Tables 3.2, 3.3 and 3.4 lists the complete encoding for the maps. Descriptions are taken from [22].

TriggerMap The TriggerMap is used for representing linedef triggers and the sectors which activates. Due to the vast amount of cases the doom engine can handle, only a few types of triggers have been considered. The mapping works by assigning an integer $i \leq 32$ to every trigger object, and subdividing triggers types in 5 groups: local doors (the ones that are activable only if the players directly interact with them), remote doors, lifts, switches and teleports. Local doors can be normal or require a key of a certain colour in order to open, but they are not indexed by the trigger index since they don't require to be linked to other linedefs in order to be opened. Table 3.5 describes the encoding for each possible item i .

RoomMap The RoomMap represent an enumeration of the rooms obtained with an algorithm that is very similar to the morphological approach used in "Room segmentation: Survey, implementation, and analysis" [3]: An euclidean distance transform [21] is first applied to the FloorMap obtaining a map that we call Distance Map or DistMap; then the local maxima are found [19] such that each maximum has a minimum distance of 3 from the closest one, resulting in the room center coordinates that are used as markers for a Watershed segmentation [18] using the negative distance map as basin. This results in a room segmentation that is good enough for descriptive purposes while maintaining good performances.

3.3.3 Graph Representation

Another way to represent the level is by using a graph. In particular, this graph is a region adjacency graph [24] built upon the RoomMap, where the nodes represent the rooms and the edges are the boundaries they have in common. This graph is built primarily for computing some features about the level and for exploiting its convenient representation of the rooms during the WAD Writing phase: this graph can be annotated with the coordinates of walls belonging to each room, and this information can be used to build a level room-by-room, with the assumption that a room could approximate a sector.

3.3.4 Text Representation

A text representation is also available, following the work of Summerville et al. in [20]. In particular the representation has been extended from one character per tile/pixel to two characters. This way it has been possible to add the information about the sector tag and the damaging floor. This representation is not currently used by our work but provided for consistence with previous works. Table 3.6 reports all the character used for this encoding.

3.3.5 Scalar Features

Each level is annotated with 176 numerical and textual features which are divided in four categories:

1. **IDGames Archive Metadata** Contain information collected from the database when levels have been downloaded. This information contains the author, the descriptions, download urls, level title, etc. Since a WAD file can contain up to 32 levels, this information is replicated for each level found in the WAD file.

Listed in table 3.7

2. **WAD-extracted features:** These features are low-level features collected directly when processing the WAD file and include the number of lines, things, sectors, vertices, the maximum and minimum coordinates, the level size in MU etc.

Listed in table 3.8

3. **PNG-extracted features** These features are computed starting from the FloorMap using an Image processing library for calculating morphological properties. Each feature is calculated both directly over the whole level and as simple statistics computed over its "floors" taken singularly. A "Floor" is intended as a part of level which is not connected to the rest of the level, thus is reachable only by means of a teleporter.

Listed in tables 3.9, 3.10

4. **Graph Features** Features computed on the room graph, inspired by the work of Luperto and Amigoni in [16]. They are used to provide a higher level representation of the level and an indicative distribution of the different room types.

Listed in table 3.11

3.4 Dataset Organization

3.4.1 Overview

This section will describe how the dataset is stored and how it can be used for future works. Since the dataset has been created primarily for instructing a neural network to generate new levels, the choices in data formats and data representation have been made to increase re-usability and flexibility in data manipulation. In particular the full dataset is stored first as a set of files indexed by a JSON dictionary, but for various reason that will be explained more in depth in chapter 6.1.4 in our work we only used a subset of levels stored as a separated archive.

3.4.2 Full Dataset and Filtered Dataset

In order to keep as much information as possible from the collected levels, we structured the representation of the DOOMDataset as follow:

- A **Full Dataset** containing all the levels we collected from the Idgames Archive.
- A set of **Filtered Datasets** containing a subset of levels that satisfy certain constraints and which are ready to use with TensorFlow [1].

Full Dataset The full dataset is kept as much portable as possible, in the sense that it shouldn't need particular technologies to be accessed except the capability of parsing JSON files

and NetworkX [13] for analysing the graph structure.

It is composed of about 9172 levels organized as a directory structure, while the maps are of different sizes given by the rescaling of the levels in MU to tile/pixel format.

The folder is structured as follow:

- **dataset.json** Contains all the scalar and textual information explained in section 3.3.5 and the relative paths to the files in sub-directories. Acts as a level database.

- **Original:** Contains all the WAD files as extracted by the archives downloaded from the IDGames Archive.

- **Processed:** Contains the Feature Maps, the graph representation, the text representation and the relative set of features.

Data in this folder is named as:

zipname_WADNAME_SLOT.json Is a json file containing all the features (3.3.5) for the level.

zipname_WADNAME_SLOT.networkx Is a gpickle compressed file containing the NetworkX graph for the level.

zipname_WADNAME_SLOT.txt Is the Text Representation (3.3.4) of the level

zipname_WADNAME_SLOT_mapname.txt Is the set of Feature Maps in PNG 8-bit greyscale format.

Where zipname, wadname, mapname indicate the name of the zip archive the WAD was stored in, the wad file itself and the Feature Map respectively, while SLOT indicates the level slot name in DOOM format (see section 3.2.2, "NAME" lump). The choice of keeping features both in the JSON database and in separated files comes from the need of recomputing the features in an easy way (for example for adding features) and for providing the possibility to manually pick or inspect level properties without the need of accessing a long json file. This, however, comes at the cost of some data redundancy.

Filtered Dataset Due to technological limitations given by the machines used for training and other reasons explained in chapter 6.1.4, data is filtered according to some criteria in order to make them uniform in term of Feature Map size and removing some level exposing extreme values of the features. Filtered dataset is stored as a set of files described as follow:

- **DATASETNAME-train.TFRecord** Dataset used for training the network using the TFRecord data format, which is a binary data format proposed by TensorFlow for improving data ingestion performances.

- **DATASETNAME-validation.TFRecord** Dataset used for model validation stored in TFRecord data format.

- **DATASETNAME.meta** Metadata and statistics about the dataset contained in train and validation datasets, in JSON format. Since the TFRecord format currently does not natively hold any information about the data contained in its records, it is useful to save data such as the item count, data structure and statistics about the dataset in a separate file, for easing the process of data normalization and other kind of operations.

In our GitHub Repository [11] it is possible to find the files relative to two different data subset:

- **128-many-floors** Dataset of images up to 128x128 pixels (smaller levels are centered and padded) which have any number of floors, consisting of 1933 levels in training set and 829 in validation set.

- **128-one-floor** Dataset of images up to 128x128 pixels (smaller levels are centered and padded) which have just one floor, consisting of 1104 levels in training set and 474 in validation set.

Provided Methods The code repository [11] hosts a Python module that provides all the necessary methods to inspect and analyse and rebuild both the full dataset and the filtered

dataset. Further information is provided in the code documentation as it goes beyond the scope of this chapter.

3.4.3 Level Size Statistics

In this section we analyse how the level dimensions in MU are distributed in the Full Dataset. Table 3.12 reports the percentile distributions for both the level height and the level width in Map units and Pixels. This distribution highlights how the choice of scaling 32 Doom units to a single pixel in order to avoid losing functional information about the levels also leads to reasonable image sizes, since more than 80 percent of the levels are representable in with an image of 256 pixels in both height and width. Joint distribution of width and height is represented in figure 3.2




Figure 3.2: Joint distribution for the features "width" (on the x axis) and "height" (on the y axis) expressed in MU for the level contained in the full dataset. It is possible to notice how the size of the majority of the levels are below 7000 Map Units

3.5 Summary

In this chapter we proposed a setting for describing the structure and the features of DOOM levels, analysing the WAD file format and exploiting its properties to extract useful data. Inspired by the work of Summerville et al. in “The VGLC: The Video Game Level Corpus” we produced a dataset of more than 9000 DOOM levels for extending the previous work and providing a ready-to-use database for future works on video-game levels. In doing so, we developed a new system for converting WAD files into a data format that is compatible with our needs, while also trying to preserve the textual format provided by the previous work for backward compatibility.

Section Length (bytes)	Section Name	Field Size	Field Name	Description
12	Header	4	Identification	ASCII string "PWAD" or "IWAD"
		4	Number of Lumps	The number of Lumps included in the WAD
		4	Table Offset	Integer pointer to the Dictionary
Variable	Lumps	-	Lump Data	Lumps stored as a stream of Bytes
16 * Number of Lumps	Directory	4	Lump Position	Integer holding a point to the lump's data
		4	Lump Size	Size of the lump in bytes
		8	Lump Name	Lump name in ASCII, up to 8 bytes long. Shorter names are null-padded

Table 3.1: WAD File structure

Value	Functional Category	Thing Description
0		Empty
1	other	Boss Brain
2	other	Deathmatch start
3	other	Player 1 start
4	other	Player 2 start
5	other	Player 3 start
6	other	Player 4 start
7	other	Spawn shooter
8	other	Spawn spot
9	other	Teleport landing
10	keys	Blue keycard
11	keys	Blue skull key
12	keys	Red keycard
13	keys	Red skull key
14	keys	Yellow keycard
15	keys	Yellow skull key
16	decorations	Bloody mess
17	decorations	Bloody mess
18	decorations	Candle
19	decorations	Dead cacodemon
20	decorations	Dead demon
21	decorations	Dead former human
22	decorations	Dead former sergeant
23	decorations	Dead imp
24	decorations	Dead lost soul (invisible)
25	decorations	Dead player
26	decorations	Hanging leg
27	decorations	Hanging pair of legs
28	decorations	Hanging victim, arms out
29	decorations	Hanging victim, one-legged
30	decorations	Hanging victim, twitching
31	decorations	Pool of blood
32	decorations	Pool of blood
33	decorations	Pool of blood and flesh
34	decorations	Pool of brains

Table 3.2: ThingsMap Encoding (1 of 3)

Value	Functional Category	Thing Description
35	obstacles	Barrel
36	obstacles	Burning barrel
37	obstacles	Burnt tree
38	obstacles	Candelabra
39	obstacles	Evil eye
40	obstacles	Five skulls "shish kebab"
41	obstacles	Floating skull
42	obstacles	Floor lamp
43	obstacles	Hanging leg
44	obstacles	Hanging pair of legs
45	obstacles	Hanging torso, brain removed
46	obstacles	Hanging torso, looking down
47	obstacles	Hanging torso, looking up
48	obstacles	Hanging torso, open skull
49	obstacles	Hanging victim, arms out
50	obstacles	Hanging victim, guts and brain removed
51	obstacles	Hanging victim, guts removed
52	obstacles	Hanging victim, one-legged
53	obstacles	Hanging victim, twitching
54	obstacles	Impaled human
55	obstacles	Large brown tree
56	obstacles	Pile of skulls and candles
57	obstacles	Short blue firestick
58	obstacles	Short green firestick
59	obstacles	Short green pillar
60	obstacles	Short green pillar with beating heart
61	obstacles	Short red firestick
62	obstacles	Short red pillar
63	obstacles	Short red pillar with skull
64	obstacles	Short techno floor lamp
65	obstacles	Skull on a pole
66	obstacles	Stalagmite
67	obstacles	Tall blue firestick
68	obstacles	Tall green firestick
69	obstacles	Tall green pillar
70	obstacles	Tall red firestick
71	obstacles	Tall red pillar
72	obstacles	Tall techno floor lamp
73	obstacles	Tall techno pillar
74	obstacles	Twitching impaled human

Table 3.3: ThingsMap Encoding (2 of 3)

Value	Functional Category	Thing Description
75	monsters	Arachnotron
76	monsters	Arch-Vile
77	monsters	Baron of Hell
78	monsters	Cacodemon
79	monsters	Chaingunner
80	monsters	Commander Keen
81	monsters	Cyberdemon
82	monsters	Demon
83	monsters	Former Human Trooper
84	monsters	Former Human Sergeant
85	monsters	Hell Knight
86	monsters	Imp
87	monsters	Lost Soul
88	monsters	Mancubus
89	monsters	Pain Elemental
90	monsters	Revenant
91	monsters	Spectre
92	monsters	Spider Mastermind
93	monsters	Wolfenstein SS
94	ammunitions	Ammo clip
95	ammunitions	Box of ammo
96	ammunitions	Box of rockets
97	ammunitions	Box of shells
98	ammunitions	Cell charge
99	ammunitions	Cell charge pack
100	ammunitions	Rocket
101	ammunitions	Shotgun shells
102	weapons	BFG 9000
103	weapons	Chaingun
104	weapons	Chainsaw
105	weapons	Plasma rifle
106	weapons	Rocket launcher
107	weapons	Shotgun
108	weapons	Super shotgun
109	powerups	Backpack
110	powerups	Blue armor
111	powerups	Green armor
112	powerups	Medikit
113	powerups	Radiation suit
114	powerups	Stimpack
115	artifacts	Berserk
116	artifacts	Computer map
117	artifacts	Health potion
118	artifacts	Invisibility
119	artifacts	Invulnerability
120	artifacts	Light amplification visor
121	artifacts	Megasphere
122	artifacts	Soul sphere
123	artifacts	Spiritual armor

Table 3.4: ThingsMap Encoding (3 of 3)

Value	Functional Category	Thing Description
0	None	Empty
10	local doors	Blue key local door
12	local doors	Red key local door
14	local doors	Yellow key local door
16	local doors	Local door
32+i	remote doors	Remote door with tag i
64+i	lifts	Lift with tag i
128+i	switch	Linedef that activates the i tag
192+i	teleports	teleport to sector i
255	exit	Level Exit

Table 3.5: TriggerMap Encoding: Each item i is connected to one or more objects. For example: switch (128+1) will open the door (32+1), raise the lift (64+1), etc.

1st character	Description	2nd Character	Description
"_"	["empty", "out of bounds"]	"_" [ascii(45)]	Empty, no tag
"X"	["solid", "wall"]	". " [ascii(46)]	Tag 1
"."	["floor", "walkable"]	"/" [ascii(47)]	Tag 2
","	["floor", "walkable", "stairs"]	"0" [ascii(48)]	Tag 3
"E"	["enemy", "walkable"]
"W"	["weapon", "walkable"]	"m" [ascii(109)]	Tag 64
"A"	["ammo", "walkable"]	"~" [ascii(126)]	Damaging floor
"H"	["health", "armor", "walkable"]		
"B"	["explosive barrel", "walkable"]		
"K"	["key", "walkable"]		
"<"	["start", "walkable"]		
"T"	["teleport", "walkable", "destination"]		
":."	["decorative", "walkable"]		
"L"	["door", "locked"]		
"t"	["teleport", "source", "activatable"]		
"+"	["door", "walkable", "activatable"]		
">"	["exit", "activatable"]		

Table 3.6: Extended Textual Representation Encoding: A second character has been added to the one used by "The VGLC: The Video Game Level Corpus": Each tile is expressed by two characters "XY" where X is the type of object and Y is the tag of the tile. Every tile that has a tag number, activates (or is activated by) the object(s) with the same tag number. So, e.g. "t/" is a teleport that leads to "T/" and "X." is a switch that activates the door "+" and possibly a floor "..."

Feature Name	Description	Type
author	Level Author	string
description	Natural language level information	string
credits	Natural language level information	string
base	Natural language level information	string
editor_used	Natural language level information	string
bugs	Natural language level information	string
build_time	Natural language level information	string
rating_value	doomworld.com level rating value	float
rating_count	doomworld.com vote count	int
page_visits	doomworld.com page visits	int
downloads	doomworld.com download count	int
creation_date	Natural language level information	string
file_url	Download page url	string
game	Doom or DoomII	string
category	doomworld.com category (eg. a-z)	string
title	Full level name	string
name	level .zip filename	string
path	relative path to wad file	string

Table 3.7: Features: IDArchive Metadata

Feature Name	Description	Type
number_of_lines	absolute number of lines in the level	int
number_of_things	absolute number of objects in the level	int
number_of_sectors	absolute number of sectors (zones with same height) in the level	int
number_of_subsectors	absolute number of subsector (convex subshapes of sectors) in the level	int
number_of_vertices	absolute number of vertices in the level	int
x_max	maximum x coordinate	int
y_max	maximum y coordinate	int
x_min	minimum x coordinate	int
y_min	minimum y coordinate	int
height	level original height in DoomUnits	int
width	level original width in DoomUnits	int
floor_height_[max min avg]	[max min avg] height for the floor	float
ceiling_height_[max min avg]	[max min avg] height for the ceiling	float
room_height_[max min avg]	[max min avg] difference between ceiling and floor height	float
sector_area_[max min avg]	[max min avg] area of sectors in squared doom map units	float
lines_per_sector[max min avg]	[max min avg] count of sector sides	float
aspect_ratio	Ratio between the longest and the shortest dimension, since a rotation of 90 of the level does not alter playability	float
walkable_area	Number of pixels the player can walk on (nonempty_size - walls)	int
walkable_percentage	Percentage of the level that is walkable	float
number_of_<things_type>	Total number of <things_type>in the level. <things_type>: { artifacts, powerups, weapons, ammunitions, keys,monsters , obstacles, decorations }	int
<things_type>_per_walkable_area	Number of <things_type>divided the walkable area in DMU.	float
start_location_[x y]_px	[x y] coordinate (in pixels, dataset format) of the start location.	int
slot	Name of the map slot. e.g "E1M1" or "MAP01"	string

Table 3.8: WAD-extracted features

Feature Name	Description	Type
floors	Number of non-connected sectors (only reachable by a teleport) of the level	int
level_area	Number of pixels composing the level	int
floors_area_[mean min max std]	[mean min max std] number of pixels composing each floor	float
level_bbox_area	Number of pixels of bounding box surrounding the level	int
level_convex_area	Number of pixels of convex hull for the whole level	int
floors_convex_area_[mean min max std]	[mean min max std] Number of pixels of convex hull for each floor	float int
level_eccentricity	Eccentricity of the ellipse that has the same second-moments as the level.	float
floors_eccentricity_[mean min max std]	Eccentricity of the ellipse that has the same second-moments as each floor	float
level_equivalent_diameter	The diameter of a circle with the same area as the level	float
floors_equivalent_diameter_[mean min max std]	[mean min max std] equivalent diameter calculated over the floors of the level	float
level_euler_number	Euler characteristic of the level. Computed as number of objects (= 1) subtracted by number of holes (8-connectivity).	int
floors_euler_number_[mean min max std]	[mean min max std] euler number over the floors of this level	float

Table 3.9: PNG-extracted features (1 of 2)

Feature Name	Description	Type
level_extent	Ratio of pixels in the level to pixels in the total bounding box. Non-empty size of the level.	float
floors_extent_[mean min max std]	[mean min max std] extent over the floors of the level	float
level_filled_area	Number of pixels of the level, obtained by filling the holes	int
floors_filled_[mean min max std]_mean	[mean min max std] filled area over the floors of the level	float
level_major_axis_length	The length of the major axis of the ellipse that has the same normalized second central moments as the level.	float
floors_major_axis_length_[mean min max std]	[mean min max std] major axis length over the floors of the level	float
level_minor_axis_length	The length of the minor axis of the ellipse that has the same normalized second central moments as the level	float
floors_minor_axis_length_[mean min max std]	[mean min max std] minor axis length over the floors of the level	float
level_orientation	Angle between the X-axis and the major axis of the ellipse that has the same second-moments as the level. Ranging from $-\pi/2$ to $\pi/2$ in counter-clockwise direction.	float
floors_orientation_[mean min max std]	[mean min max std] orientation over the floors of the level	float
level_perimeter	Perimeter the level which approximates the contour as a line through the centers of border pixels using a 4-connectivity.	float
floors_perimeter_[mean min max std]	[mean min max std] perimeter over the floors of the level	float
level_solidity	Ratio of pixels in the level to pixels of the convex hull image.	float
floors_solidity_[mean min max std]	[mean min max std] solidity over the floors of the level	float
level_hu_moment_[0 ... 6]	Hu moments (translation, scale and rotation invariant).	float
level_centroid_x	Centroid coordinate x	float
level_centroid_y	Centroid coordinate y	float

Table 3.10: PNG-extracted features (2 of 2)

Feature Name	Description	Type
art-points	Number of articulation points in the room adjacency graph. An articulation point is a node which removal would result in a bipartite graph	int
assortativity-mean	Mean assortativity. Assortativity is the tendency of one node to be connected with similar nodes.	float
betw-cen-[min max mean var]	Node centrality statistic calculated with the betweenness method	float
betw-cen-[skew kurt]	Node centrality statistic calculated with the betweenness method	float
betw-cen-[Q1 Q2 Q3]	Node centrality statistic calculated with the betweenness method	float
closn-cen-[min max mean var]	Node centrality statistic calculated with the Closeness method	float
closn-cen-[skew kurt]	Node centrality statistic calculated with the Closeness method	float
closn-cen-[Q1 Q2 Q3]	Node centrality statistic calculated with the Closeness method	float
distmap-[min max mean]	Maximum value in the distance map, i.e. the size of the largest room. Background is ignored from computation.	float
distmap-[var skew kurt]	Mean value for the distance map, i.e. the mean room size. Background is ignored from computation.	float
distmap-[Q1 Q2 Q3]	Skewness of the distnace distribuion Background is ignored from computation.	float

Table 3.11: Graph features

Percentile	Width in DU	Width in pixels	Height in DU	Height in pixels
10th	2305	72	2065	64
20th	3017	94	2657	83
30th	3633	113	3137	98
40th	4161	130	3606	112
50th	4737	148	4101	128
60th	5393	168	4657	145
70th	6177	193	5313	166
80th	7212	225	6209	194
90th	9040	282	7809	244
100th	31233	976	32743	1023

Table 3.12: Percentiles of level width and height distributions in both Map Units and Pixels

Chapter 4

System Design and Overview

Overview This chapter describes the proposed system from an high level perspective. The purpose of this chapter is in fact to give an overview of how the system modules interact from the point of view of the data flow and use case analysis. Section 4.3 describes the system from the perspective of a Data Flow analysis, highlighting the major transformations that produce new levels. Section 4.2 introduces the main system components from a use case perspective, while section 4.1 focuses on the GAN structure. Details of the chosen architecture are left to Chapter 6.1.1.

4.1 Generative Model Structure

Overview The initial system design was built upon the architecture of Generative Adversarial Networks [12] from Goodfellow et al. Given the problem of generating video-games levels, the need of having a way to control to some extent the generation process naturally arose. For this reason, we adopted a conditional version [17] of the GAN Model, proposed by Mirza and Osindero, applied to a more recent GAN model that is discussed in section 6.1.1.

Conditional GAN Structure We present in figure 4.1 the general architecture which defines the inputs and the outputs of the generative model we are using. This structure refers to the Conditional Generative Neural Network we introduced in chapter 2.1.1. In particular, figure 4.1 shows the general working principles of a GAN: It is composed by two neural networks, namely a Generator and a Discriminator (or Critic, depending on the underlying architecture that is adopted).

The input of this subsystem are defined as follows:

- X : Batch of images having m channels, corresponding to the Feature Maps.
- Y : Batch of vectors having one component for each scalar feature considered.
- Z : Batch of random noise vector, typically sampled from a Uniform or Gaussian distribution.

The discriminator network takes as input a vector of images X and a vector of features Y , while the generator network G takes as input a the vector of features Y and a vector of random noise Z , which is used like a "random seed" for generating different samples from the data distribution. We use the subscript "True" or "Gen" to distinguish from the Feature Maps coming from the input dataset and those that are generated by the generator G .

For what concerns the network outputs, we have that:

$X_{Gen} = G(Z|Y)$: Samples generated from the generator network

$Logits(X_{Gen}) = D(X_{Gen}|Y)$: Discriminator output when real samples are in input to the discriminator.

$Logits(X_{True}) = D(X_{True}|Y)$: Discriminator output with generated samples are provided to

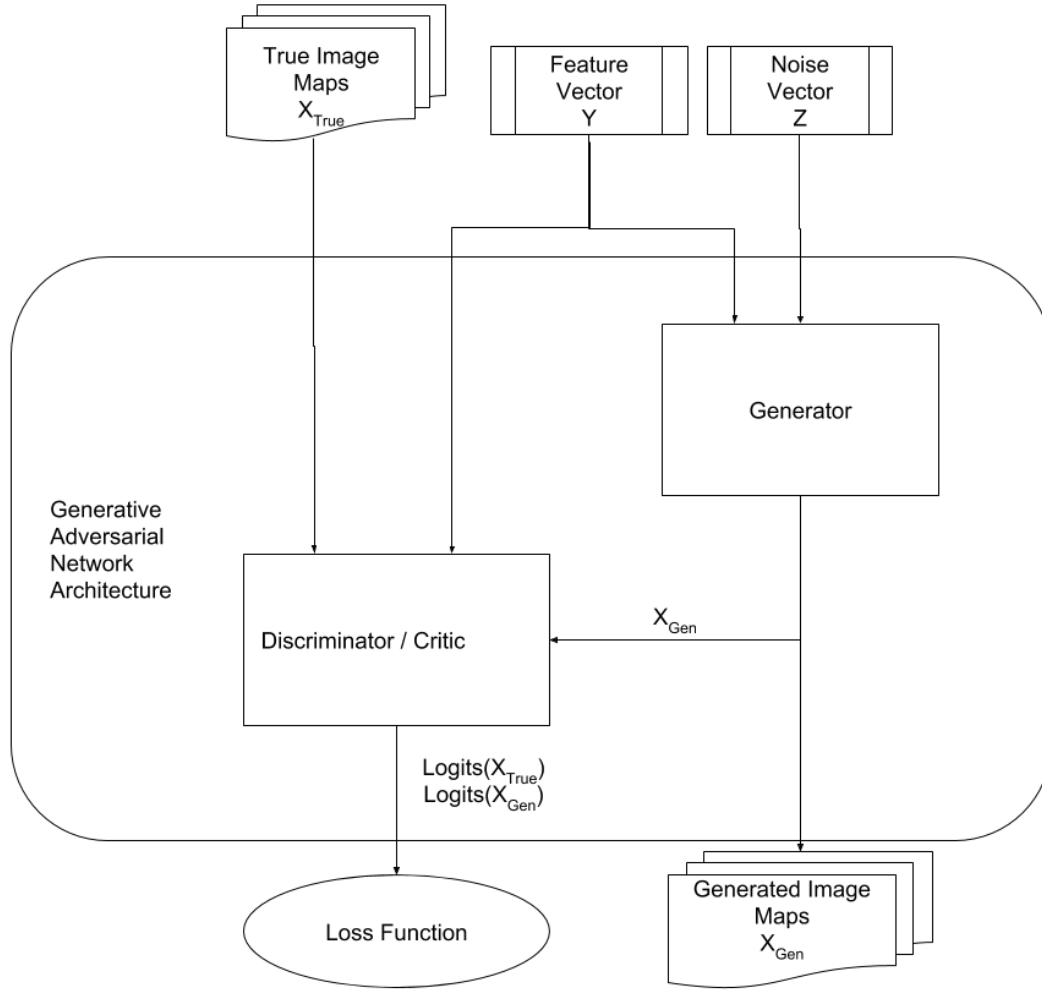


Figure 4.1: System Overview: Generative Model Structure. A Conditional GAN is composed of a generator model and a discriminative model. The discriminative model takes as input either the Images coming from the dataset or the ones generated by the generative model. Both networks are conditioned by the Y feature vector, while the generator also takes an input a noise vector Z to sample from the data distribution it is approximating.

the discriminator.

Logits are actually the output of the last layer of the discriminator network before the last *activation function*. A Loss function for either the generator and the discriminator is written upon those values and they are alternatively optimized to train the entire network. All the details are given when we'll describe the chosen GAN Architecture in section 6.1.1.

4.2 Use Cases

Overview This section will describe the main use cases of our system, which are useful for replicating our results. The emphasis is put on how inputs and outputs are used in each case, while the internal structure of the generative model is not represented to simplify the notation. Every figure in this section also describes the function of the dataset metadata introduced in section 3.4.2 while describing the filtered dataset. In particular the neural network inputs and

outputs are limited in a certain range, typically between 0 and 1 or -1 and 1. For this reason, dataset statistics are needed to properly rescale input and output data. Blocks which are written in bold type indicate those inputs and outputs that are of interest for the corresponding use case.

4.2.1 Use Case: Model Optimization (Training)

This use case describes the optimization of the model, which is commonly referred as the training phase. This is the phase in which data from the dataset is fed into the model and the loss functions are minimized in order to find the network weights that allow the generation of samples. In the ideal case the generated samples should come from a distribution which is undistinguishable from the true data distribution. In reality, this is limited by both the ability of the discriminator in selecting the features that allow to distinguish true data from artificial one, and the generator ability is misleading the discriminator in its task. Figure 4.2 shows that in this use case both the Scalar Features Y and the Feature Maps Y from the dataset are used to train the network. At each epoch the generator network is fed with a random vector Z along with the conditioning vector Y . This procedure produces a training loss that is provided to an optimizer that optimizes the network parameters.

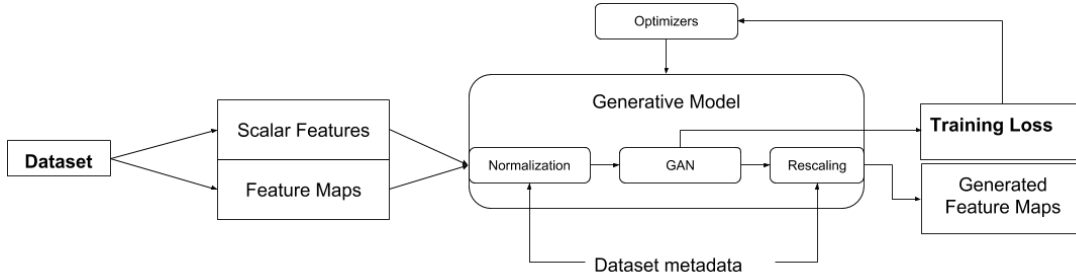


Figure 4.2: Use Case: Model Optimization.

4.2.2 Use Case: Model Validation and Sample Evaluation

This use case is used to assess the ability of the model to generate new samples on previously unseen feature vectors. This is accomplished by running two different procedures in which only the validation set, composed of samples that are left out from the training phase, is used:

1. A *validation loss* is calculated by feeding the discriminator with images $X_{True,Val}$ coming from the validation set and their corresponding feature vector Y_{Val} . This approach is often used in standard (non generative) neural networks where it is a good measure of network generalization capabilities. As discussed in section 6.1.1 in our setting, however, this may not always be a meaningful metric with every proposed underlying architecture. This is usually due to the fact that with many GAN architectures the loss does not correlate well with the quality sample.

2. A set of *quality metrics* (6.1.6) are computed directly on the true samples $X_{True,Val}$ and the samples $X_{Gen,Val} = G(Z|Y_{Val})$ generated by conditioning the network with the same features of $X_{True,Val}$. This is based on the assumption that if the network actually learns a correlation between the Y vector of features and a certain set of features proper to the corresponding X samples, then the true sample and the generated one might show a certain grade of similarity according to the given features. Since this may be a strong assumption, especially in the setting where we are not able to measure what features are actually mapped to each component Y_i due to the high dimensionality of the problem, a set of more generic metrics are chosen and presented in section (6.1.6). Selected metrics should be general enough to express, when averaged on batches of samples, a concept of "sample quality" without directly referring to the features encoded in the Y vector.

Figure 4.3 shows how the scalar features from the validation set are used to generate new samples, that are compared to the corresponding true images.

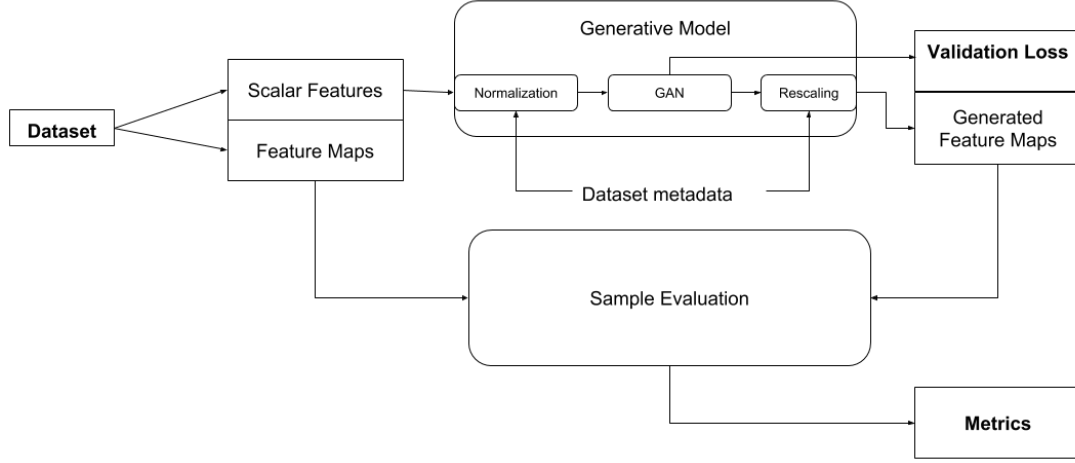


Figure 4.3: Use Case: Validation and Sample Evaluation.

4.2.3 Use Case: Sampling or Generation

This use case is the one that produces the final output and it is run when a final model is selected. In particular, our system supports several methods for sampling the network in order to generate new levels.

These different sampling methods are motivated in section 6.2, while in this section we only highlight the differences from an input/output point of view. The main problem of sampling this network is choosing a feature vector that has enough probability in the distribution

- **Dataset Sampling:**
- **"Factors" Sampling:**
- **Direct Sampling:**
- **"Content" Sampling:**

4.3 Data Flow

Introduction

Description Figure 4.5 represent the main data transformations for training a new model to produce new DOOM Levels starting from a set of levels collected from the IDGames Archive. Since the scope of this character is to only show a general overview of the system, not every component and transformation is described in detail. The figure is organized as a process developing from left to right. In particular, blocks on the left represent the inputs to the process and blocks on the right boundary are the produced artefacts.

It is possible to identify three main data paths: The first produces the model parameters and it is identified with the use case "Model Optimization", or "Training" (4.2.1), in which the Training Set is used to find a set of network weights that allow to produce new levels.

The second one corresponds to the Model Validation and Sample Evaluation use case (4.2.2) in which the model is used to generate new levels using the feature vector taken from the validation set, then the samples from the validation




Figure 4.4: Use Case: Sampling or Generation.

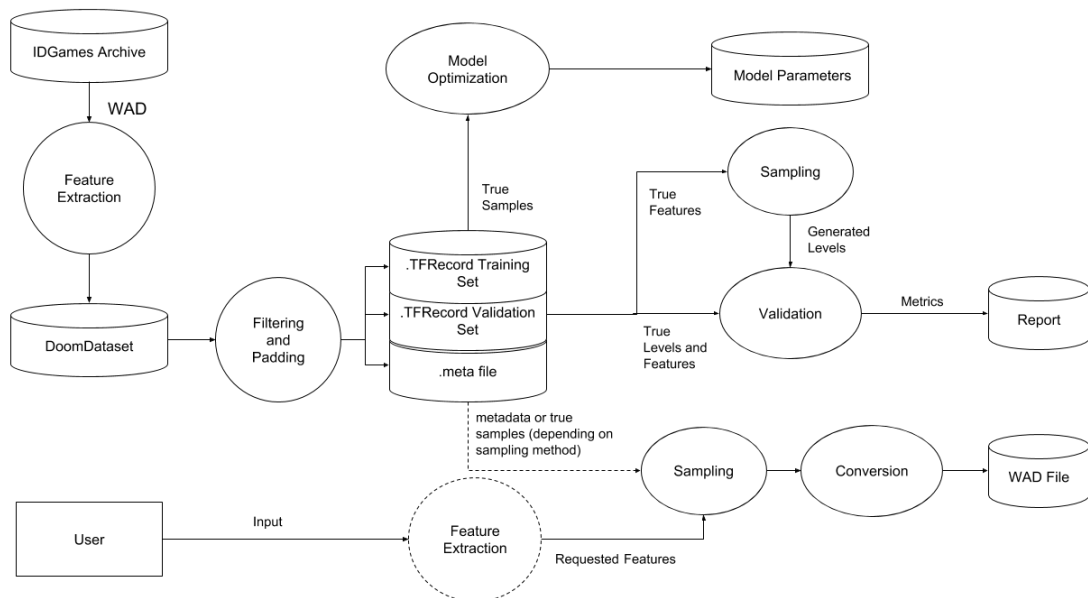


Figure 4.5: System Overview: Data Flow Diagram. Disc-shaped blocks represent data archives, circles represent processes and transformations, while the labels on arrows represents intermediate data. Dashed lines represents objects which presence depends on the use case of reference.

4.4 WAD Editor and Feature Extractor

4.5 Summary

Chapter 5

System Architecture

5.1 Component View

5.2 Neural Network Architecture

Chapter 6

Experiment Design and Results

6.1 Neural Network Architecture and Training Algorithm

6.1.1 Network Architecture and Hyper-parameters

Chosen Architecture

Advantages

6.1.2 Techniques and "GAN Tricks" used

6.1.3 Training Algorithm

6.1.4 Input data

6.1.5 Resulting Model

6.1.6 Evaluation metrics

6.2 Sampling the network

6.3 Generated Samples

6.4 In-Game Demonstration

6.5 Summary

Chapter 7

Results Evaluation and Conclusions

7.1 Results Evaluation

7.1.1 Samples Evaluation

7.1.2 Loss of accuracy

7.2 Summary

Chapter 8

Future Work

8.1 Open Problems

8.2 Possible Applications and future develops

Chapter 9

Appendix

Glossary

deathmatch Game mode in which two or more players compete against each other in achieving the highest number of kills before a timeout or a player reaches a predetermined score.. 55

DU Doom Units, see MU. 55

Feature Map Image describing a particular aspect of a level. 26, 29, 55

floor Unconnected piece of level, reachable only by means of a teleporter. 55

FloorMap Image describing which parts of the map are occupied by a floor and which are empty.. 26–28, 55

GAN Generative Adversarial Network. 55

HeightMap Image describing the floor height of a level. 27, 55

Linedef Line that connects two vertices in the WAD file specification. 24, 26, 55

Lump Any section of data within a WAD file.. 22, 24, 55

MU Map Units, Coordinate unit used in Doom Rendering Engine.. 23, 24, 26, 28–30, 55, 57

RoomMap Image describing the room segmentation of the level. 27, 55

Sector A Sector in a DOOM level is any closed area (with possibly invisible walls) that has a constant floor and ceiling height and texture. plural. 23, 55

Sidedef A "vertical plane" in the WAD file specification.. 24, 55

Thing Any deployable asset of a DOOM level, such as Enemies, Power-Ups, Weapons, Decorations, Spawners, Etc. plural. 23, 55

ThingsMap Image describing how the game objects are placed inside the level. 27, 55

TriggerMap Image describing doors, lifts and switches and their correlation. 27, 55

WAD Default format of package files for the DOOM / DOOM II video-games. "WAD" is an acronym for "Where's all the Data?" [23]. 21–23, 26, 29, 55, 57

WallMap Image describing the impassable walls in a level. 26, 55

Bibliography

- [1] Martín Abadi et al. *TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems*. Software available from tensorflow.org. 2015. URL: <https://www.tensorflow.org/>.
- [2] Barry Bloom. *Barry Bloom post on OFC.UNT.EDU News*. URL: <https://groups.google.com/d/msg/alt.games.doom/Wjp0gf-zBf8/YIYsTQyD9F4J>.
- [3] R. Bormann et al. “Room segmentation: Survey, implementation, and analysis”. In: *2016 IEEE International Conference on Robotics and Automation (ICRA)*. 2016, pp. 1019–1026. DOI: [10.1109/ICRA.2016.7487234](https://doi.org/10.1109/ICRA.2016.7487234).
- [4] Colin Phipps, Simon Howard, Colin Reed, Lee Killough. *BSP v5.2*. [BSP - The Doom node builder software, accessed November 2017]. 1994 - 2006.
- [5] Doomworld on DoomWiki. *Doomworld - The Doom Wiki at DoomWiki.org*. [Online; accessed 06/02/2018]. 2005. URL: <https://doomwiki.org/wiki/Doomworld>.
- [6] Doomworld.com Website. *Doomworld*. [Online; accessed 16/10/2017]. 1998. URL: <https://www.doomworld.com/>.
- [7] Encyclopedia of Mathematics. *Orientation*. URL: <http://www.encyclopediaofmath.org/index.php?title=Orientation&oldid=39859>.
- [8] *Encyclopedia of Optimization*. Kluwer, 2001. ISBN: 9780792369325. URL: <https://books.google.com.mx/books?id=gtoTkL7heSOC>.
- [9] Matthew S Fell. *The Unofficial Doom Specs*. Online. 1994. URL: <http://www.gamers.org/dhs/helpdocs/dmsp1666.html>.
- [10] Fuchs, Henry and Kedem, Zvi M. and Naylor, Bruce F. “On Visible Surface Generation by a Priori Tree Structures”. In: *SIGGRAPH Comput. Graph.* 14.3 (July 1980), pp. 124–133. ISSN: 0097-8930. DOI: [10.1145/965105.807481](https://doi.org/10.1145/965105.807481). URL: <http://doi.acm.org/10.1145/965105.807481>.
- [11] Edoardo Giacomello. *DoomPCGML*. <https://github.com/edoardogiacomello/DoomPCGML>. 2017-2018.
- [12] Ian J. Goodfellow et al. “Generative Adversarial Networks”. In: *CoRR* abs/1406.2661 (2014). arXiv: [1406.2661](https://arxiv.org/abs/1406.2661). URL: <http://arxiv.org/abs/1406.2661>.
- [13] Aric A. Hagberg, Daniel A. Schult, and Pieter J. Swart. “Exploring network structure, dynamics, and function using NetworkX”. In: *Proceedings of the 7th Python in Science Conference (SciPy2008)*. Pasadena, CA USA, Aug. 2008, pp. 11–15.
- [14] John Carmack. *Doom Engine Source code on GitHub.com*. [Online; accessed 06/02/2018]. 1997. URL: <https://github.com/id-Software/DOOM>.
- [15] Y. Lecun et al. “Gradient-based learning applied to document recognition”. In: *Proceedings of the IEEE* 86.11 (1998), pp. 2278–2324. ISSN: 0018-9219. DOI: [10.1109/5.726791](https://doi.org/10.1109/5.726791).
- [16] Matteo Luperto and Francesco Amigoni. “Predicting the Global Structure of Indoor Environments: A Constructive Machine Learning Approach”. unpublished article at the moment of writing. 2018.

- [17] Mehdi Mirza and Simon Osindero. “Conditional Generative Adversarial Nets”. In: *CoRR* abs/1411.1784 (2014). arXiv: [1411.1784](https://arxiv.org/abs/1411.1784). URL: <http://arxiv.org/abs/1411.1784>.
- [18] P. Neubert and P. Protzel. “Compact Watershed and Preemptive SLIC: On Improving Trade-offs of Superpixel Segmentation Algorithms”. In: *2014 22nd International Conference on Pattern Recognition*. 2014, pp. 996–1001. DOI: [10.1109/ICPR.2014.181](https://doi.org/10.1109/ICPR.2014.181).
- [19] scikit-image development team. *Scikit Image Documentation - peak_local_max Documentation*. URL: http://scikit-image.org/docs/0.12.x/api/skimimage.feature.html#skimimage.feature.peak_local_max.
- [20] Adam James Summerville et al. “The VGLC: The Video Game Level Corpus”. In: *Proceedings of the 7th Workshop on Procedural Content Generation* (2016).
- [21] The Scipy community. *SciPy Documentation - distance_transform_edt*. 2008-2009. URL: https://docs.scipy.org/doc/scipy-0.14.0/reference/generated/scipy.ndimage.morphology.distance_transform_edt.html.
- [22] Thing Types on DoomWiki. *Thing Types - The Doom Wiki at DoomWiki.org*. [Online; accessed November 2017]. 2005. URL: http://doom.wikia.com/wiki/Thing_types.
- [23] Tom Hall. *DOOM Bible, Appendix A: Glossary*. [Online; Doomworld Mirror, accessed 06/02/2018]. 1992. URL: <http://5years.doomworld.com/doombible/appendices.shtml>.
- [24] Alain Trmeau and Philippe Colantoni. “Regions Adjacency Graph Applied To Color Image Segmentation”. In: *IEEE Transactions on Image Processing* 9 (2000), pp. 735–744.